Deploying a spam message detection application using R over Docker and Kubernetes

SAGAR VORA[1,*] **AND RAHUL SINGH**[1]

[1]*School of Informatics and Computing, Bloomington, IN 47408, U.S.A.*
[*]*Corresponding authors: vorasagar7@gmail.com, rahul_singh919@yahoo.com*

In the last few decades, online spam has become one of the major problem for the sustainability of the Internet. Due to the excessive amount of spams, the quality of information available on the Internet has reduce drastically. Moreover spam messages are also creating problems among the various search engines available and the web users. This report aims at developing an application which would detect spam messages from actual meaningful messages using Pandas and R. For the purpose for parallelizing the process, we would deploy the application using Docker containers on the Kubernetes cluster using ansible scripts which would automate the deployment.

Keywords: Docker, Ansible, Kubernetes, R, Pandas, Spam, <add spam detection algorithms>

https://github.com/cloudmesh/sp17-i524/raw/master/project/S17-IR-P007/report/report.pdf

---

INTRODUCTION

Today, the Internet [1] has been adopted rapidly in the day to day life of people. It has provided a platform for information generation and consumption. Moreover, it is used on a daily basis to search for information and acquire knowledge. The online encyclopedia Wikipedia™ [2] provides a good example of a more socialized Internet because the content within Wikipedia™ is collectively generated by its users, rather than webmasters or designated editors. The ease with which content can be generated and published has also made it easier to create spam. Spam can be stated as any information which does not add value to a user of the web. Messages which are inappropriate, unsolicited, repeated and irrelevant can be all classified as spam.

So in this report, we are provinding an application that would identify valid messages and spam messages from a given dataset. For spam detection, we are using various techniques like Bayesian, <text here> ...... Moreover, deploying our application using Docker [3] containers on the Kubernetes [4] cluster will give a distributed approach.This would also speed up the process of identifying the spam messages. We have also deployed it on different cloud environments like Chameleon, JetStream, FutureSystem and have performed benchmark analysis of the application. This would let us the time taken by the algorithm on these cloud solutions.

Name of the Technology	Purpose in the Project
R	data analytics
Docker	container for the applicat...
Kubernetes	cluster creation and manage...
Cloudmesh Client	An client application used to ssh in various...
Ansible	Automation language to deploy application...

Fig. 1. Technologies used in the Project

SOFTWARE STACK

ELEMENTS OF THE PROJECT

Ansible

No one likes repetitive tasks, so with Ansible [5], IT admins can begin automating away the drudgery from their daily routine tasks. Ansible is a simple automation language that can perfectly describe an IT application infrastructure. Ansible is an open source automation engine which can be used to automate cloud provisioning, configuration management, and application deployment. It can also perform more advanced IT tasks such as continuous deployment or rolling out updates with zero

downtime.

A major difference in Ansible and many other tools in the space is its architecture.

Architecture

Ansible is an agentless tool,it doesn't requires any software to be installed on the remote machines to make them manageable. By default is manages remote machines over SSH or WinRM, which are natively present on those platforms [**?**].

Like the other configuration management software, Ansible distinguishes between two types of servers: one being the controlling machines and other being the nodes. Ansible uses a single controlling machine where the orchestration begins. Nodes are then controlled by a controlling machine over SSH [**?**]. The location of the nodes are described by the inventory of the controlling machine.

Ansible modules are deployed by Ansible over SSH. These modules are temporarily stored in the nodes and communicate with the controlling machine through a JSON protocol over the standard output

Playbooks

Playbooks [6] are Ansible's configuration, deployment, and orchestration language. They let us control the remote systems with a policy which we might want them to enforce. If Ansible modules act as tools in your workshop, then playbooks are your instruction manuals, and your inventory of hosts are your raw material. Playbooks can be used to manage configurations of and deployments to remote machines. They can sequence multi-tier rollouts involving rolling updates, and can delegate actions to other hosts, interacting with monitoring servers and load balancers.

Ansible Galaxy

R

R is a language and environment for statistical computing and graphics [7]. Pandas does not provide a significant statistical modeling environment as it is still a work in progress. R provides a variety of statistical model analysis, classification, clustering and graphical techniques to provide this environment. Integrating Python's efficiency with R's capability allows us to build a highly a desirable analysis model for our application.

Docker

Docker is an open-source project that automates application deployment by packaging the application in *containers*. Containers provide application portability by bundling together an application and its needed resources in a package so that they can be deployed on different platforms without worrying about resource dependencies. Application containerization is an OS (Operating System) level emphvirtualization for deploying and running an application instance without launching a virtual mahine for each application [8]. A container has its own environment variables, filesystem and libraries that is needed by the application, thus eliminating OS or hardware dependency. Containers abstract the OS kernel while a *VM (Virtual Machine)* hypervisor abstracts an entire device.

Docker allows application developers to package their applications into isolated containers. Docker automates the repetitive tasks of setting up and configuring development environments thus allowing developers to focus only on building software. A dockerized application can simply ship between platforms as the complexity of software dependencies is handled by the container.

Docker standardizes container creation and can be used to pack, ship and run an application as a lightweight container that can run in any environment. Docker can be integrated with other devops applications like Puppet, Chef, Vagrant, Ansible and Kubernetes. We shall use Docker with Ansible and Kubernetes in our project.

Dockerfile and DockerImage

To package an application and its dependencies in a single file, docker introduces the concept of a *docker image*. The docker engine creates a docker image by parsing contents of a *dockerfile*. A dockerfile is a script composed of various commands to build a container in a step-by-step, layer-by-layer manner [9]. Once an image has been built it can be shared with other users by pushing it to a public repository on *DockerHub* or *GoogleCloudPlatform*. In this manner, an image once built by the docker engine can be used across the organization by making a docker pull request.

Kubernetes

Kubernetes is an open-source platform for automating deployment, management and scaling of containerized applications across a cluster [10]. The more granular an application is, the more components it consists of and thus requires management of these components. Kubernetes helps in faster deployment of applications and scaling them on the fly. Moreover it optimizes the use of hardware by using the resources which are needed. Kubernetes provides container management features like component replication, load-balancing, service-discovery and logging across components [11]. A Kubernetes cluster can be deployed on either physical or virtual machines. We shall be deploying a kubernetes cluster using *kubeadm* - the kubernetes command line tool and *Minikube* which is a lightweight Kubernetes implementation which creates a VM on the local machine and deploys a simple cluster containing only one node. The Minikube CLI provides basic bootstrapping operations for working with the cluster, including start, stop, status, and delete commands.

Kubernetes Terminologies

Kubernetes defines the following set of primitives which provide mechanisms for deploying and scaling applications.

Pods

A pod is the smallest unit of a kubernetes cluster and has a unique ip-address within the cluster. A pod consists of one or more containers that can share resources and can be controlled as a single application [10] [12]. Thus all the involved containers in a pod are scheduled on the same host. A pod can be thought of as a single virtual machine in terms of resource sharing and scheduling. Pods can be managed manually using the *Kubernetes API* or can be managed by a controller.

Services

A kubernetes service is a collection of pods that perform the same function and are presented as a single entity. This way a service can be emphasized as a one tier of a multi-tier application. Service act as an interface to a group of containers so that service-consumers need only reference the single access location. Kubernetes facilitates service discovery by assigning a stable IP address and a namespace to a service. This idea abstracts the change of IP addresses of pods within a service that result due to pod failure or pod rescheduling.

Replication Controllers

A replication controller is a framework for horizontal scaling of pods. Semantics of a pod are defined in a <pod_name>.yaml file which also defines the replication details that need to be done. The replication controller performs replication by scaling a number of pods across a cluster based on the pod definition file [10]. The replication controller has to make sure that a certain number of copies of a pod are always up and running. Thus, in event of a pod failure it replaces the failed pod with a new replica.

Labels and Selectors

Kubernetes allows users or internal components to assign a key-value pair tag to any API object in the system. An object can have one or more labels associated with it, but each with a unique key. eg: appversion = 1.0, development_stage = 4. Label selectors are queries against labels that return matching objects [12]. This way each object in the system can be referenced with a single label or a combination of multiple labels for fine grained control.

Kubernetes architecture

Kubernetes exhibits the master-slave architecture. The components can be split into those that manage an individual node(slave) and those that manage the master or control plane.
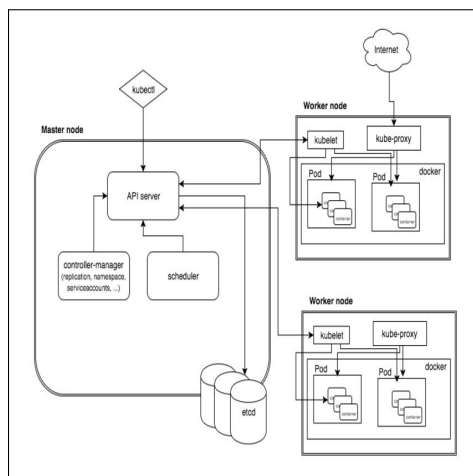

Fig. 2. Kubernetes architecure [11]

Kubernetes Master Node/Control plane

The Kubernetes master node is responsible for managing the kubernetes cluster and orchestrating the worker nodes, where the actual pods are scheduled. The master node can work as a single node cluster where it can schedule pods to work on the master node itself by tainting the schedule policy rule on the node. The master node also referred as the control plane consists of several components:

API Server

The API server is the most fundamental component of the Kubernetes master and serves as the entry point for all the REST commands used to control the cluster. The API server serves up the Kubernetes API using JSON over HTTP, providing both internal and external interface to the cluster [12] [13]. It validates the REST requests, executes them and updates the status of the objects in the *etcd* storage.

etcd storage

etcd is a simple, distributed and consistent key-value store that stores configuration data of the cluster and represents the state of the cluster at any point of time [10]. Kubernetes uses etcd for service discovery and provides a simple HTTP/JSON API as an interface for setting or retrieving values from the store. Other components watch the state of the etcd store to bring themselves up to the desired state. Data being stored in the etcd store are deployed services, pods, replication information etc.

Scheduler

The scheduler component is responsible for the deployment of pods and services on the cluster nodes. The scheduler has the information about the availability of resources on a node and schedules unscheduled pods on the nodes accordingly. Along with scheduling, the scheduler also tracks resource utilization of each node and ensures that workload scheduled is not in excess to the resources available [10].

Controller-manager

The controller manager is the process embedding the different types of controllers like the Replication Controller or the DaemonSet Controller on a kubernetes master. The controllers query the API Server to manipulate the resources like pods,services etc. which they manage.

Worker Node

The worker node also called as minion node is where the containers are actually deployed. The worker contains all the necessary services needed to manage the networking between containers, communicate with the master node and assign resources to the scheduled containers [11]. Every worker node must run the container runtime i.e docker and other components stated below to ensure proper communication with the master.

Docker

Docker runs on each of the worker nodes. It is responsible for downloading the docker images and running the configured pods by starting the container.

Kubelet

Kubelet gets the pod definition from the api-server and is responsible for maintaining the pod in the desired state. Kubelet is the worker service that monitors the health of each pod and communicates the status of each node via a heartbeat message to the master. If the pod is not in the desired state, it is redeployed to the same node [10]. Kubelet is also responsible for communicating with the etcd storage to get information about the services and update the storage about newly created ones.

Kube-Proxy

Kube-proxy acts as a network proxy and a load balancer. It is responsible for networking of TCP and UDP packets to the appropriate container based on the IP address of each packet [10] [11].

Kubectl

kubectl is a command line tool that communicated with the API server to fetch important information about the nodes, pods, services and events in the cluster.

DESIGN

Building the Classification model

CrossValidation for the training data

To address the problem of incoming spam messages, a model shall be developed using the Bayesian Classification technique to correctly classify each incoming email/text message as a spam or a legitimate one. The model aims at developing a message filter that shall correctly classify messages based on word probabilities that are extracted from the training dataset.The training dataset to build the model consists of 5574 message records. Dataset taken from [14]. The training process shall use the cross-validation feature provided by R to build the classification model and use Bayes theorem of conditional probability to predict the class of each incoming message.

To develop an efficient training model, we shall partition the data into 2 subsets - training data and classification data. We shall choose one of the subsets for training and other for testing. In the next iteration the roles of the subsets shall be reversed, i.e the training data becomes the classification one and vice versa. This operation shall be carried out until each individual record is used both as a classification and training record. We shall use the cross validation feature provided by R for this subsampling. This subsampling technique handles the underfitting problem and guarantees an effective classification model.

Training process

Content of each of the spam marked messages shall be processed through Naive Bayes Classifier. The classifier shall maintain a bag of words along with the count of each word occuring in the spam messages. This word count shall be used to calculate and store the word probability in a table that shall be cross-referenced to determine the class of the record on classification data [15].

A selected few words have more probability of occuring in a spam messages than in the legitimate ones. Eg: The word "Lottery" shall be encountered more often in a spam message. The classifier shall correlate the bag of words with spam and non-spam messages and then use Bayes Theorem to calculate a probability score that shall indicate whether a message is a spam or not. The results shall be verified with the results available on the training dataset and the classifier accuracy shall be calculated. The classifier shall use the Bayesian theorem over the training dataset to calculate probabilities of such words that occur more often in spam messages and later use a summation of scores of the occurence of these word probabilities to estimate whether a message shall be classified as spam or not. After working on several samples of the training dataset, the classifier shall have learned a high probability for spam based words whereas, words in legitimate message like family member or friends names shall have a very low probability of occurence.

Classifying new data

Once the training process has been completed, the posterior probability for all the words in the new input email is computed using Bayes theorem. A threshold value shall be defined to classify a message into either class. A message's spam probability is computed over all words in its body and if the sum total of the probabilities exceeds the predefined threshold, the filter shall mark the message as a spam [16].

A higher filtering accuracy shall be achieved through filtering by looking at the message header i.e the sender's number/name. Thereby if a message from a particular sender is repeatedly marked as spam by the user, the classifier need not evaluate the message body if it is from the same sender.

DEPLOYMENT

Our application will be deployed using Ansible [5] playbook. Automated deployment should happen on two or more nodes clouds or on multiple clusters of a single cloud. Deployment script should install all necessary software along with the project code to Kubernetes cluster nodes using the Docker image.

Deployment process - Ansible

<TODO - describe the ansible deployment and scripts>

Dockerizing the application

To containerize our application, we need to create a docker image for it. To create a docker image we need to create a docker file that describes the semantics to create the docker image. Docker version used for this process is 1.12.6.

The Dockerfile

Docker can build images automatically by reading instructions from a Dockerfile. A dockerfile is a series of text commands that a user can run on a command line to build an image [?]. Instructions within a Dockerfile have the following format:

```
INSTRUCTION   arguments
```

Fig. 3. Syntax of every command in a dockerfile

Docker reads the instructions from a dockerfile in the order they are written. The very first line of a docker file must be a 'FROM' clause which specifies the base image from which the image is being built. As our application works with R, we need to specify the base package as one that shall help us run our R application. *DockerHub* has a r-base package that binds the latest version of R and its libraries together that we shall use [?]. The argument following the FROM clause is a repository/tag name that the docker engine automatically looks up on http://hub.docker.com.

Following the FROM statement, we can specify all actions we need to perform like specifying the work directory for building the image or copying files into our work directory so that every resource is available under a single folder. Our application needs several text mining packages to be available as a part of the installation and thus we use the docker RUN command to download these dependencies and make them part of the docker image. The R library in the docker container shall pre-install all the dependencies as the image is built and we shall only care about running the application logic. Lastly, our docker file ends with a CMD statement that specifies the command that the container is supposed to execute on successful instantiation. CMD accepts an array of command names followed by the parameters. Since we want our image to execute our R application we shall specify the command as stated in the dockerfile below.

```
FROM r-base
COPY . ~/myscripts
WORKDIR ~/myscripts
RUN Rscript -e "install.packages('tm')"
RUN Rscript -e "install.packages('e1071')"
RUN Rscript -e "install.packages('RWeka')"
RUN Rscript -e "install.packages('ada')"
RUN Rscript -e "install.packages('rbenchmark')"
CMD ["Rscript","spamdetection.r"]
```

Fig. 4. Dockerfile of the application

Creating a docker image

Once the Dockerfile is built, we can build the image using docker build command as follows :

```
cc@rahpsing-055:~/dockerDirectory$ sudo docker build
-t  spamdetectionapplication  ~/dockerDirectory/

Sending build context to Docker daemon 490.5 kB
Step 1 : FROM r-base
latest: Pulling from library/r-base

5d797a4cfc3e: Already exists
19c2f092956c: Already exists
93a6e2a5f4fa: Already exists
8f6ce7170232: Already exists
d1474396bacc: Already exists
e729120880f2: Already exists
Digest: sha256:e192edf861d61caff0b329436...
Status: Downloaded newer image for r-base:latest
 ---> 16fe32463daa

Step 2 : COPY . ~/dockerDirectory/
 ---> a6717da2ec47
Removing intermediate container 8fb4aaaec0d8

Step 3 : WORKDIR ~/dockerDirectory/
 ---> Running in 0220b193ed11
 ---> a3c4d57e23c9
Removing intermediate container 0220b193ed11

Step 4 : RUN Rscript -e "install.packages('tm')
      .
      . <installation_logs>
      .
Step 9 : CMD Rscript spamdetection.r
 ---> Running in 4edb44f2ffb6
 ---> a4814eeadf35
Removing intermediate container 4edb44f2ffb6
Successfully built a4814eeadf35
```

Fig. 5. Building an image line by line from a dockerfile

The built image is placed in our machine's local docker registry and can be viewed with the 'docker images' command.

```
$ docker images
REPOSITORY                 TAG      IMAGE ID     SIZE
spamdetectionapplication   latest   a4814eeadf   646.1 MB
r-base                     latest   16fe32463d   645.6 MB
```

Fig. 6. Listing images in the local repository

Running the application

We shall run our application using the docker run command.

```
docker run spamdetection
```

Fig. 7. Running a docker image

We can check the container id of our application along with other important information using the 'docker ps' command.

Sharing the docker image

By default, the docker CLI points to docker's public registry which is located at http://hub.docker.com. We need to create a docker account to upload our application image so that it could be directly referenced by kubernetes later. Once a docker account is registered, create a public repository and give it a name. The repository shall be accessible by the format '<username>/repositoryname'.

To tag the current docker image that we created on our system we need to inform the docker engine to point to our registry. To do so, we shall use login command provided by the docker CLI.

```
docker login
```

Fig. 8. Linking a dockerhub account to the local registry

The login request shall ask for a valid docker account userid and password. Once authenticated, the docker engine shall establish a link with the registry associated with the username. To upload our docker image to the repository, we need to tag the image. Docker provides the 'tag' command to do so.

```
docker tag image_name <username>/<repository_name>:
<tag>
```

Fig. 9. Tagging a local docker image

The <tag> can be any name given by the user to uniquely identify the image. Once tagged, we need to push the image to the repository using the command line 'docker push' command.

```
docker push rahpsing/kubernetesi524:spamdetection
```

Fig. 10. Pushing an docker image to a registry

As the image is now publicly available on the repository, we can do a simple 'docker pull' from any machine to deploy the container with no dependencies. This is possible as the image wraps together the application code file, the data file and dependent r-base package from the latest version of R available on dockerhub's library.

Deployment via Minikube
TODO

Deployment via kubeadm
Kubeadm is a part of kubernetes 1.4 distribution that allows users to install and set up a Kubernetes cluster. Kubeadm works with local VM's, cloud servers or physical servers [17].

Creating a cluster
To create a kubernetes cluster, it is essential to install docker, kubelet, kubectl and kubeadm on all the machines that are to be a part of the network. Kubelet shall help to start pods and containers on all the machines of the network. kubectl allows us to monitor the activities of the cluster once it is up and running. Primarily, it is used only on the master node. Kubeadm is used to setup the cluster by allowing multiple worker nodes to bind with the master on a unique network identifier token.

Once the mentioned components are installed, we need to initialize the master so that it can accept requests from other service nodes to join the cluster. Master initialization is done with the 'init' command.

```
kubeadm init
```

Fig. 11. The master initialization command

Executing the above command tells kubernetes that the host machine shall serve as the master in the cluster. Kubeadm initializes all the other dependent components like the API server and generates a unique key that identifies the master in the network. Client nodes can issue a join request with the master by using the unique identifier as a part of the join command.

```
kubeadm join <unique_token>
```

Fig. 12. Syntax of the join command

A few seconds after running the join command we can query the master to list the available nodes in the cluster using 'kubectl get nodes' command.

```
kubeadm join --token 87ce11.50ab6a5eea 192.168.0.204
$ kubectl get nodes
NAME            STATUS    AGE
rahpsing-056    Ready     2m
rahpsing-057    Ready     2m
```

Fig. 13. Registering nodes to the cluster

Before the master is ready to schedule pods, it is imperative that the API server and the kubedns service are up and running. We can check the status of all the services of the system using the '–all-namespaces' agrument along with the "get nodes" command.

The kubedns service is responsible for networking with its client nodes. Without the kubedns service, the kubernetes master shall be unable to schedule pods on the worker nodes as pods cannot communicate with each other. To get the service up and running, we need to install a pod network. Kubernetes provides many addon services that can be used to setup the cluster network policy and enable networking. Of the list of available addons we shall use 'weave-net' as it provides us with a network policy and does not require an external database [?].

```
kubectl apply -f https://git.io/weave-kube
```

Fig. 14. Creating the weave-net pod-network

```
$ kubectl get pods --all-namespaces
NAME                                    READY    STATUS
etcd-rahpsing-056                       1/1      Running
kube-apiserver-rahpsing-056             1/1      Running
kube-controller-manager-rahpsing-056    1/1      Running
kube-discovery-2849056221-5xws4         1/1      Running
kube-dns-2247936740-m9198               3/3      Running
kube-proxy-amd64-d91cy                  1/1      Running
kube-proxy-amd64-ycg5e                  1/1      Running
kube-scheduler-rahpsing-056             1/1      Running
weave-net-1sk03                         2/2      Running
weave-net-p9cpd                         2/2      Running
```

Fig. 15. Services running in the cluster

Creating a pod
A pod as described initially is the smallest unit of work in a kubernetes cluster. Semantics of a pod are defined in a pod.yaml file which is then passed to the kubectl CLI to initialize the pod and maintain its desired state.

The pod.yaml file specifies the containers that compose a pod. Each container works on top of a docker image that it runs once instantiated. We shall use the spamdetection image for our container that we created by dockering our application and pushing the image to the repository. By default, kubernetes searches for image tags specified in the pod.yaml file on docker's public repository. If the matching image is found, the pod is created successfully or it results in a pod failure.

Once the semantics of a pod are fixed we can use kubectl to instantiate the pod. We can also use 'kubectly apply' if we wish to make changes to the definition of the pod and want the implementation to reflect it.

```
kubectl create -f pod.yaml
```

Fig. 16. Creating a pod

kubectl provides 'get XXX' commands to query status of any type of object in the system. We can check the status of our pods by using the API's get pods command.

```
$ kubectl get pods
NAME                READY    STATUS     AGE
spamdetectionimage  0/1      Running    46s
```

Fig. 17. Status of pods in the cluster

To get a detailed view of each element we can use the 'describe <element_name>' command where <element_name> is unique under each type of component in the system.

Creating a deployment

Similar to creating a pod, a deployment can be created in Kubernetes using the same yaml file syntax. The only notable difference in this case is the value of the 'type' and the 'apiVersion' field in the yaml file. The 'type' field is set to 'deployment' in case of creating a deployment while the apiVersion field is set to 'extension/v1beta1' as version 'v1' which was used for pod creation does not support deployments. With the deployment file, we have the flexibility of specifying the number of replicas we want to deploy. We set the number of replicas to 2, which means that kubernetes will ensure that at any given point in time the system will have 2 replicas of our application running. After defining the number, we define the containers whose replicas we wish to maintain.

```
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: spamdetectionapplication
spec:
  replicas: 2
  template:
    metadata:
      labels:
        app: echo
    spec:
      containers:
      - name: container2
        image: rahpsing/kubernetesi524:spamdetection
            application
        ports:
        - containerPort: 80
```

Fig. 18. Deployment file of the application

kubectl command to create a deployment is the same as the one we used to create a pod. The Kube Control API parses the yaml file and understands the type of object it is expected to create.

As with pods, the command to list the deployments running on the cluster is 'kubectl get deployments'. Now that our deployment is created we can check its status by using various kubectl commands. As we set the number of replicas to be 2, there should be atleast 2 pods running our application which can be verified as below.

```
$ kubectl get pods
NAME                         READY    STATUS
spamdetectionapplication-1   1/1      Running
spamdetectionapplication-2   1/1      Running
```

Fig. 19. Deployment with 2 replicas

Fetching the application output

As stated previously, a pod is a collection of multiple related containers. Our application executes in of the containers within the pod. Since our application prints output to the console, we shall be able to view it by accessing logs of the relevant pod.

```
kubectl get pods | grep spamdetectionapplication*

kubectl logs <pod_name>
```

Fig. 20. Fetching logs for the application

As each pod is hosted on a separate node with a unique ip address, we can issue a curl or a wget request to the node to fetch the output. Since our application output spans over multiple lines we shall redirect and save the output to a file so that it can be referenced at any point of time with ease.

```
curl https://<node_ip>:<node_port>/containerLogs/
default/<pod_name>/spamdetectionapplication
--insecure > output.txt
```

Fig. 21. Saving container logs to a file

Following the above steps, the application output shall be available in the file 'output.txt' in the current directory.

Scheduling pods on the master

By default, kubernetes does not allow pods to be scheduled on the master node for security purposes. However, this setting can be overridden by removing the taint on the master node and allowing pod execution.

```
kubectl taint nodes --all node-role.kubernetes.io/
master-
```

Fig. 22. Command to schedule pods on the master

Execution of the above statement shall remove the taint 'node-role.kubernetes.io/master' from all the nodes in the network including the master node and thus allowing the scheduler to schedule pods across the network [17]. This feature also allows us to create a single node cluster.

ISSUES TACKLED WITH KUBERNETES SETUP AND DEPLOYMENT

ImagePullBackOff

```
$ kubectl get pods
NAME              READY    STATUS          RESTARTS
myapplication     0/1      ImagePullBackOff 21
```

Fig. 23. The ImagePullBackOff error

The ImagePullBackOff error indicates that the container was unable to fetch the desired image from the repository specified. This happens in case of improper tag names associated with the application image or reference to a repository that doesnt exist. By default, any text specified against the 'image' key in the pod or deployment file makes a direct lookup on dockerhub. Thus any docker image that has been built on the system cannot be referenced by the kubernetes engine unless it is tagged an uploaded to a public repository on dockerhub. Dockerhub allows users to create public repositories for no charges. Thus, the only way to use your custom created image is to register for an account on dockerhub, create a public repository, upload your local image with a unique tag name and use the repository and tag name in the yaml file. This issue took a lot of our development time as there is no clear documentation which explicitly describes

the kubernetes interpretation of the value of the 'image' key specified in the pod yaml file.

CrashLoopBackOff

```
$ kubectl get pods
NAME              READY  STATUS           RESTARTS
myapplication   0/1     CrashLoopBackOff  49
```

Fig. 24. The CrashLoopBackOff error

Kubernetes allows us to execute instructions after a container has been deployed successfully. The output of these instructions is visible in the container's logs which can be accessed by using the 'get logs' primitive of the kubectl CLI. Commands to be executed can be specified under the 'args' key of the container specification in the pod definition file. However, kubernetes defines strict parsing for those arguments with no pointers to indicate where the execution failed. Any issues with the yaml file shall trigger a crashloopbackoff error with no traces of line numbers or clear log messages to guess what went wrong. It was after considerable research and trial and error that we figured out that an extra space between arguments to the sleep command was causing the error.

Kube DNS service not working

```
$ kubectl get pods --all-namespaces
NAME                                   READY  STATUS
etcd-rahpsing-056                       1/1     Running
kube-apiserver-rahpsing-056             1/1     Running
kube-controller-manager-rahpsing-056    1/1     Running
kube-discovery-2849056221-5xws4         1/1     Running
kube-dns-2247936740-m9198               0/3     Container
                                                Creating
```

Fig. 25. Failure of the kube-dns service

Kubedns is an integral part of the kubernetes master setup as it enables pods to communicate with each other. Before creating any pods or deployments we need to setup a pod-network. The pod network defines security policies and establishes networking within the cluster. Kubernetes provides various addons to be used as its pod network. We tried to use *flannel* but were unable to bind it to our network and get the dns service working. After numerous attempts we figured out that *weave-net* is the best pod-network that we could use as its network was updated to work with the CNI networking policy introduced in Kubernetes 1.4. Also, using weave net as our pod network allows us to use the weave-kube service where we can observe our cluster components and their configurations on a GUI accessible by a web browser.

Unsuccessful download requests from within a kubernetes pod

Initially, our R application program started with many download and install-package statements as we needed several text mining packages to implement the respective algorithms. Even after successful deployment, we couldn't find the application output in the container logs. It took us a fair amount of time to figure out that something wasn't working correctly as the pod logs failed to highlight the issue. To debug in deeply, we started deploying applications with simple print statements to check their behavior. After several attempts, we could draw a conclusion that deployment would result in a failure if the program included package import statements that needed to be downloaded from an external repository. We ran the same application image with docker only to find it working perfectly. Thus, we could state that, from within a kubernetes pod, an application is unable to reach the network and download the desired packages. We could find no precise explanation of this behavior on the kubernetes forum or issues opened on stack overflow or kubernetes git page.

As a workaround, we segregated the download statements from the R program and included it within our docker file. This modification would download and install the R libraries and packages as a part of the build process of the docker image, eliminating the need to install it post deployment. With the dependent libraries as a part of the docker container, the application program would run successfully and generate the output as a part of the logs for the pod it is associated with.

DISCUSSION

TBD

CONCLUSION

TBD

APPENDICES

TBD

REFERENCES

[1] "What is internet?" Web Page, accessed: 2017-04-12. [Online]. Available: http://searchwindevelopment.techtarget.com/definition/Internet

[2] "Wikipedia," Web Page, accessed: 2017-04-12. [Online]. Available: https://www.wikipedia.org/

[3] "What is docker?" Web Page, accessed: 2017-04-02. [Online]. Available: https://www.docker.com/what-docker

[4] "Kubernetes," Web Page, accessed: 2017-03-10. [Online]. Available: https://kubernetes.io/docs/concepts/overview/what-is-kubernetes/

[5] "Ansible, deploy apps. manage systems. crush complexity," Web Page, accessed: 2017-04-12. [Online]. Available: https://www.ansible.com/it-automation

[6] "Playbook," Web Page, accessed: 2017-04-12. [Online]. Available: http://docs.ansible.com/ansible/playbooks.html

[7] "R:what is R?" Web Page, accessed: 2017-04-02. [Online]. Available: https://www.r-project.org/about.html

[8] M. Rouse, "What is application containerization (app containerization)?" Web Page, accessed: 2017-04-18. [Online]. Available: http://searchitoperations.techtarget.com/definition/application-containerization-app-containerization

[9] O.S.Tezer, "Docker explained: Using dockerfiles to automate building of images," Web Page, Dec 2013, accessed: 2017-04-24. [Online]. Available: https://www.digitalocean.com/community/tutorials/docker-explained-using-dockerfiles-to-automate-building-of-images

[10] Wikipedia, "Kubernetes-wikipedia," Web Page, accessed: 2017-04-18. [Online]. Available: https://en.wikipedia.org/wiki/Kubernetes

[11] N. Isabekyan, "Introduction to kubernetes architecture," Web Page, Jul 2016, accessed: 2017-04-24. [Online]. Available: https://x-team.com/blog/introduction-kubernetes-architecture/

[12] J. Ellingwood, "An introduction to Kubernetes," Web Page, accessed: 2017-04-18. [Online]. Available: https://www.digitalocean.com/community/tutorials/an-introduction-to-kubernetes

[13] K. Marhubi, "Kubernetes from the ground up:the API server," Web Page, accessed: 2017-04-24. [Online]. Available: http://kamalmarhubi.com/blog/2015/09/06/kubernetes-from-the-ground-up-the-api-server/

[14] "SMS spam collection dataset," Web Page, accessed: 2017-03-10. [Online]. Available: https://www.kaggle.com/uciml/sms-spam-collection-dataset

[15] J. Provost, "Naive-Bayes vs. Rule-Learning in Classification of Email," in *Artificial Intelligence Lab*. The University Of Texas at Austin: The University Of Texas, 1999, accessed: 2017-03-10. [Online]. Available: http://mathcs.wilkes.edu/~kapolka/cs340/provost-ai-tr-99-281.pdf

[16] Wikipedia, "Naive bayes spam filtering," Web Page, January 2017, accessed: 2017-03-10. [Online]. Available: https://en.wikipedia.org/wiki/Naive_Bayes_spam_filtering

[17] "Using kubeadm to create a cluster|kubernetes," Web Page, accessed: 2017-04-24. [Online]. Available: https://kubernetes.io/docs/setup/independent/create-cluster-kubeadm/